

Performance Analysis of an Offline Digital Euro Prototype

Robbert Koning, Johan Pouwelse (thesis supervisor)

Distributed Systems
Delft University of Technology
Delft, The Netherlands

—master’s thesis—

Abstract—This document is a model and instructions for \LaTeX . This and the `IEEEtran.cls` file define the components of your paper [title, text, heads, etc.]. *CRITICAL: Do Not Use Symbols, Special Characters, Footnotes, or Math in Paper Title or Abstract.

Index Terms—component, formatting, style, styling, insert

I. INTRODUCTION

In recent years, the European Central Bank (ECB) has been exploring the possibility of realizing its own Central Bank Digital Currency (CBDC), the ‘digital Euro’. The ECB has published various reports and resources that outline the need for such a project (i.e. [1], [2]). Calls for expression of interest are being published and the ECB aims to complete its investigation phase by October 2023 [3], [4]. The main reason for this development is the rise of digital payments and corresponding decline of cash usage. According to reports published by De Nederlandsche Bank (DNB), the national bank of the Netherlands, the share of cash payments dropped from 56% in 2010 to 21% in 2020 [5], [6]. The Swedish Riksbank mentions similar trends for Sweden [7].

Euro cash is the only public form of money that is directly backed by the ECB [2]. Digital payments are not; they are backed by private parties such as commercial banks. A critical dependence on these parties is eroding the sovereignty of the Euro. They cannot safeguard reliability comparable to that of ECB-backed cash. Nevertheless, there is demand for reliability, especially in times of crisis [8]. In recent history, there have been several financial crises that caused large-scale bankruptcies which consequently impacted consumer’s savings (e.g. in 2008). CBDCs can provide reliability and safeguard consumers against the effects of large-scale bankruptcy of commercial payment providers.

Foreign organisations, commercial parties, and cryptocurrencies are threatening the influence of central banks. A report published by the ECB discusses the risk of *currency substitution*. Substitution occurs when a new form of money, unregulated by ECB, gains major usage in the EU. The new payment method would likely have to outperform its competitors, for instance by being cheaper and/or more convenient. According to the report, currency substitution could have a

range of negative effects on the ECB’s monetary policy and even threaten the EU’s independence [1]. The actors responsible for this consternation are mostly large corporations and foreign central banks [2], [9]. Some interested governmental parties are e.g. the United States government and the People’s Bank of China [10], [11]. An interested commercial party is for instance Meta (formerly Facebook), which initiated *Diem*¹, a hypothesized stablecoin that did not launch due to legal and regulatory issues. Due to the potentially far-reaching impact of the introduction of CBDCs to consumers and the rapid pace with which central banks are operating, public discourse on the topic has been increasing. Some regulators are questioning the use case of CBDCs and their implications on the role of central banks and consumer’s privacy [12], [13].

To determine the appropriate role for a CBDC in the EU, and to compete with other payment solutions, the ECB has launched an extensive exploratory phase. In this exploration, the ECB has expressed interest for its CBDC to be usable in an offline environment [1]. This is crucial in case of network failure or in areas without a reliable internet connection. A prominent example of currency that is spent offline is cash.

This thesis concerns itself with implementing a simple transferable digital currency on the Kotlin-IPv8² protocol stack and doing a performance analysis. In accordance with the *Offline First* design principles, the currency can be spent offline and guarantees retroactive fraud detection. It is therefore resilient against temporary failure of central servers, unlike many currently deployed systems. This thesis contributes 1) a software-implemented simple token-based transaction system 2) a performance analysis of various bottlenecks in this system, with a special emphasis on the Kotlin-IPv8 framework and 3) a slightly optimized version of the EVA² protocol.

II. PROBLEM DESCRIPTION

The main difficulty with implementing offline digital currency is the *double spending problem*. Double spending is the action of spending a digital unit of value more than once, illegitimately. In a digital environment this is possible because

¹For Diem, refer to <https://www.diem.com/en-us/>

²For Kotlin-IPv8 and EVA, refer to <https://github.com/Tribler/kotlin-ipv8>.

currency is easily duplicated. This makes fraud prevention difficult, especially in offline scenarios. In such scenarios, verifying transactions is hard due to limited communication.

The double spending problem has never been solved in an offline setting, only in an online setting. Many cryptocurrencies (e.g. Bitcoin) mitigate the problem by utilizing ‘global consensus’ [14]. This removes the need for a central authority but does require near-immediate connectivity to parts of the network. Global consensus disallows offline transfers and is therefore not a well-suited solution to make offline spending possible.

The concept of digital currency is not new; it is widely agreed upon that the idea was first proposed by Chaum in 1983 (see Section III-A) [15]. Since then, the problem of offline spending has been explored extensively. From the literature we observe that most currency schemes are *token-based*, as opposed to *account-based*. Token-based schemes transfer tokens; monetary units that can be identified with a serial number. By contrast, account-based schemes perform monetary transfers by crediting and debiting accounts. The crucial difference is that currency in token-based schemes is identifiable, whereas in account-based schemes it is not. A commonly used analogy is that token-based schemes are comparable to banknotes, whereas account-based schemes are comparable to bank deposits. A crucial lesson observed from the literature and our main prior work (see Section III-B), is that account-based systems complicate robustness measures such as safeguarding against double spending [16].

We believe that a token-based scheme lends itself better for offline spending. A token-based scheme requires the generation of tokens—analogue to minting coins—. The token minting process and transaction protocol of our implementation are described in Section IV. We limit ourselves to an implementation where currency is represented by digital tokens of a fixed value.

An additional lesson from the literature is that robust realisations are lacking for numerous theoretical proposals made over the last 39 years. The realisation of many of the difficult designs is by itself a difficult challenge. Another lesson from the literature is that adoption requires public support. This is out of scope for the efforts presented in this thesis.

III. RELATED WORK

A. Advancements in digital currency

In 1983, Chaum introduced blind signatures in what is widely accredited as the first paper to describe digital currency [15]. The paper describes a novel cryptographic primitive, the ‘blind signature’. It allows parties to sign messages without knowing their contents. The result is that the signing party cannot relate their own signature to the original message they signed. With this primitive, the literature’s first digital cash scheme was described. In this scheme, a monetary authority guarantees the validity of payments. Due to blind signatures, the authority cannot identify the recipient of any transaction it verifies, thereby safeguarding consumers’ privacy. Chaum’s cash was however *non-transferable*. Non-transferable e-cash

can be spent only once, after which it must be redeemed by a trusted authority. The authority returns an equivalent amount of cash that is spendable again.

In 1989, Okamoto introduced *transferable* e-cash [17]. Transferable e-cash is more like physical cash; it can be spent repeatedly, from one user to another. It does not require a network connection to a monetary authority with every transaction. In the same paper, *divisible* e-cash was introduced. In contrast to physical cash, divisible e-cash can be spent in smaller denominations than the piece that is owned. An advantage of divisible e-cash is that exact payments can be made and change is not required.

In 1995, a modification to blind signatures was proposed that made them ‘fair’ [18]. Most blind signature schemes were *perfectly unlinkable*. Perfect unlinkability means that no monetary authority can relate withdrawals to payments. Therefore, these schemes allowed for a variety of crimes to be undetectable, such as money laundering. With the introduction of ‘fair’ blind signatures, an additional and independent authority (such as a judge) would be able to obtain information that can be used to detect crime.

In 2008, Bitcoin was presented, widely accredited as the first major cryptocurrency. It solves the double spending problem probabilistically and without a central authority [14]. Bitcoin’s value is determined by market forces and highly volatile. This is in stark contrast to CBDCs, which are tethered in value to government-issued money.

B. Eurotoken

We consider the main prior work for this thesis to be the first Eurotoken prototype by Delft University of Technology [19]. This digital currency is also implemented on Kotlin-IPv8. Eurotoken is an account-based system and is non-transferable by default. Eurotoken opted for a trusted authority to verify transactions. Likewise, it is therefore not decentralized. The advantage of this approach in the context of CBDCs is that it enables the respective central bank to exert control over the network. Moreover, it provides a non-deterministic near-immediate transaction finality.

Based upon Eurotoken and in line with many proposed digital cash schemes, we also sacrificed decentralization and opted for a monetary authority. By contrast, our prototype is token-based and offline transferable by design.

C. Price Stability

It is fundamental for a European CBDC to be tethered in value to the Euro. A high price volatility like Bitcoin’s is undesirable for a medium of exchange [20]. There are various ways in which the value of an asset can be kept stable. This topic has gained renewed interest with the rise of ‘stablecoins’—cryptocurrencies that aim to be non-volatile with regards to a major non-cryptocurrency or physical asset. There is an inverse relationship between the potential stability of stablecoins and how much they are decentralized [21]. The strongest stabilization mechanism is collateralization by currency or off-chain assets such as gold. By allowing free

trade between a stablecoin and its collateral at a fixed price, arbitrage prevents the stablecoin’s price from fluctuating greatly. However, off-chain assets are not traded in a decentralized way and as such there is a trade-off between decentralization and stability. To the best of our knowledge, no decentralized and highly stable stablecoins exist. The prototype described in this thesis makes use of an implied centralized exchange. The implementation of this or other means of maintaining price stability is intentionally left out of scope.

IV. DESIGN AND ARCHITECTURE

This research implements a centralized CBDC prototype that allows offline transactions with fixed-value tokens and guarantees retroactive fraud detection.

The proposed system requires a trusted monetary authority that is in charge of token exchange and transaction verification. We refer to this party as ‘authority’ and identify them by their public key. Verification is therefore a centralized operation. The motivation for this design choice is elaborated upon in Section III-B. The process of exchanging currency for tokens is beyond the scope of this thesis and is briefly discussed in Section VII. Our design is deliberately simple to make it robust and well understood.

All system participants apart from the authority are clients. They, too, are identified by their public key. It is assumed that clients know the public key of the authority in the network. It is also assumed that authorities know the real identities of clients. While this is not necessary for the proposed system to function, implicating a public key with fraud loses its severity if the instigator can remain anonymous. This is discussed further in Section IV-H.

Clients can transact tokens to each other and consult the authority to verify the validity of their tokens. If clients cannot connect to the authority, for instance during a power outage, they can continue transacting but defer verification until they can connect.

To realize retroactive fraud detection, the implemented system requires authorities to be able to unambiguously reconstruct the sequence of owners of a token. This is done by providing each token with a linked list of all previous owners until its last verification. Details of this procedure are explained further in this section.

A. Token Format

The token protocol is based upon transacting tokens. A diagram of a token is given in Figure 1. Each token contains:

- 1) *Serial number*. An 8-byte unique token identifier.
- 2) *Value*. A 1-byte representation of the token’s worth. Like cash, tokens have a limited number of fixed denominations. Certain byte values are mapped to certain denominations; the remaining values are considered invalid.
- 3) *Authority public key*. A 74-byte public key³ of the authority that is in charge of the token (the ‘authority’).

³Public keys in Kotlin-IPv8 are 74 bytes long: 10 bytes for a string prefix; 32 bytes for an encryption key; and 32 bytes for a verification key. Only the latter is required for our prototype. However, Kotlin-IPv8 does not allow these to be split by design, as parties are identified by the entire 74 bytes.

1st Recipient owns:

Serial No. Value Nonce	Authority's signature
------------------------------	--------------------------

2nd Recipient owns:

Serial No. Value Nonce	Authority's signature	1st Recipient's signature
------------------------------	--------------------------	------------------------------

Nth Recipient owns:

Serial No. Value Nonce	Authority's signature	1st Recipient's signature	...	N-1th Recipient's signature
------------------------------	--------------------------	------------------------------	-----	--------------------------------

Fig. 1. Graphical representation of a token. Tokens represent monetary units of fixed value that store all their previous recipients until they are verified by an authority.

- 4) *Nonce*. A 64-byte pseudo-random nonce used by the authority to differentiate between differing occasions where the same token is sent to the same recipient.
- 5) *Recipients*. A list of recipient-proof pairs in chronological order. This list must contain at least a first pair:
 - a) *First recipient public key*. A 74-byte public key³ of the token’s first recipient after creation or validation.
 - b) *First proof*. A 64-byte signature (‘proof’) given by the authority signing *Serial number*, *Value*, *Nonce*, and *First recipient public key*.

All pairs in the list are of the same format and bit-length. The second pair (if present) contains *Second recipient public key* and a signature given by *First recipient public key* signing *First proof* and *Second recipient public key* together. Likewise, all subsequent pairs follow the same pattern; they contain a signature by the previous public key in the list, signing the previous proof together with the next public key. This signature chain corresponds to the token changing ownership during transactions.

The bit-lengths of the signatures and public keys were adapted from those used in Kotlin-IPv8 and are not integral to the protocol’s functioning.

The initial size of a token when transferred from a monetary authority to the first recipient adds up to 285 bytes. Every additional recipient adds $74 + 64 = 138$ bytes for its public key and signature. When a token is verified by a monetary authority, its size is reset to 285 bytes (see Section IV-E).

B. Token Minting

When a token is created, its *Serial number*, *Value*, *Nonce*, *Authority public key*, and *Recipients* list are set as specified in Section IV-A. The authority stores a copy of the entire token and sends it to the intended client.

C. Client Verification

When a client obtains a token, it verifies it in a 3-step process. First, the client verifies that the token’s last recipient

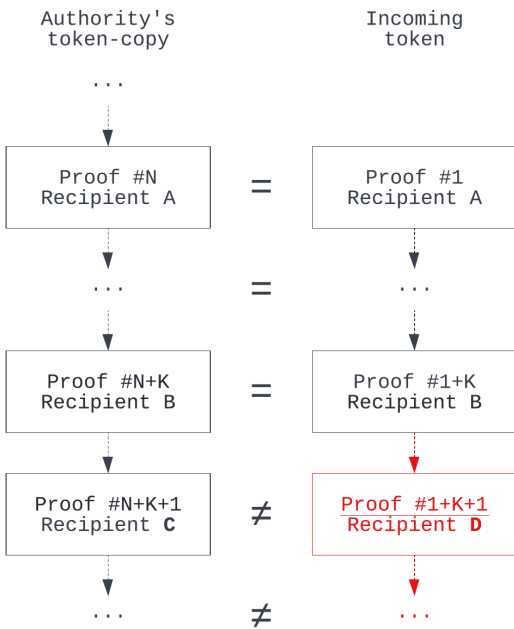


Fig. 2. The authority's double spending detection mechanism. In the figure, recipient B double spent a token, which was detected because proof $N+K+1$ of the authority was not equal to proof $1+K+1$ of the incoming token.

(that is, the last public key in the *Recipients* list) refers to them. Second, the client verifies that it knows the token's *Authority public key* and that this key created the token's *First proof*. Third, the client verifies the remaining chain of proofs in the *Recipients* list. The purpose of the client's verification process is merely to ensure that they have received an unambiguous proof of transfer from their transaction's counterparty. This proof can later be used by the relevant authority to proof potential fraud. A client deciding that a token is valid does not imply that an authority will decide the same. The client's verification does however guarantee that clients victimized by fraud can proof so eventually.

D. Client Transaction

A token's initial recipient may choose to send it to another client. If it does, it must append a new pair to the token's *Recipients* list that contains the desired recipient's public key and a signature of the token's last proof together with the desired recipient's public key. This is depicted in Figure 1.

E. Authority Verification

The authority's verification process is started when a client sends them a token to verify. The verification process contains 6 steps:

- 1) The authority ensures that the received token has more than 1 recipient in its *Recipients* list. If not, the token is either invalid or ineligible for verification.
- 2) The authority ensures that the token's last recipient is the client that sent the token in for verification.
- 3) The authority queries if the token is still valid. The knowledge that the authority once signed the received

token, which can be derived from the token's *First proof*, says little about the token's current state. The authority compares its public key against the token's *Authority public key* and queries the token's *Serial number* to ensure that itself is the authority that manages the token. Then it verifies that the token is still in circulation and not e.g. blacklisted.

- 4) The authority will, like an honest client, verify the chain of proofs in the *Recipients* list.
- 5) The authority will attempt to detect double spending by comparing the proof of the last pair ('last proof') of its token-copy to *First proof* in the received token. If these are identical, double spending cannot be proven (see Section IV-F) and the authority will finalize verification. Finalizing verification requires the authority to update its copy of the token by appending all new recipient-proof pairs of the received token to its *Recipients* list. It will also append a new pair containing the desired recipient—the one who sent the token for verification—and a corresponding proof.
- 6) The authority copies the verified token, empties the *Recipients* list save for its last entry, and sends the verified token to the desired recipient.

F. Double Spending Detection

In Section IV-E it is mentioned that the authority updates its token-copy's *Recipients* list upon a valid verification. This means that its last proof is updated as well. To detect double spending, an authority compares the last proof of its token-copy to *First proof* in the received token. A diagram of this scenario is depicted in Figure 2.

If a token is double spent, then multiple versions of the token will eventually reach their authority. The first time, double spending cannot be detected and the token-copy is updated. Subsequent times, the authority's token-copy already has an updated *Recipients* list and therefore its last proof does not correspond to the double spent token's *First proof* anymore. Thus, double spending must have occurred if the proofs differ. If the proofs are equal, double spending might have occurred.

When double spending is detected, the authority will search for the instigator. It will find the received token's *First proof* in the *Recipients* list of its token-copy. It will then compare the recipient-proof pairs of the token-copy with those of the received token. Comparison starts from the pairs that contain *First proof*. All pairs before it have already been verified. Eventually, it must find two differing pairs, after which all pairs will be different because proofs are chained to each other. The first differing pairs are the start of the token's split history and proof that double spending was performed by the client that signed them.

G. Replay Attack Prevention

The detection mechanism of Section IV-F allows for a replay attack in an offline environment. If a malicious sender *A* were to replay sending the same token to the same receiver

B as before, said receiver would not flag this as malicious behavior. If B in turn were to spend this token, upon verification of the token, B would be flagged as a double spender. When an authority compares the transaction history of the token, it cannot distinguish A 's first transaction to B from its second. Thus B spending the token is the first occurrence that differs from the authority's history. As described in Section IV-F, B is therefore marked as a fraudster.

There exist various solutions for preventing such an attack. One such solution is to initiate a transaction with the receiver sending a short handshake that includes a pseudo-random nonce. The sender must include this nonce in its transaction to proof with overwhelming probability that they did not replay the transaction. Another solution is to have receivers maintain a list of the last proofs of all tokens they have ever received.

H. Anonymity

For offline usage, the implemented system requires aggregating a linked list of previous owners of a token, up until the last verification by an authority. Specifically, recipients of a token can see all previous recipients of that token until its last verification. This is detrimental to privacy and anonymity. There are digital cash schemes that provide stronger notions of anonymity. Some schemes protect the identities of previous recipients and provide *unlinkability*, such that it is also impossible to relate different payments from the same client [22]. Some schemes provide an even stronger notion of anonymity where an adversary cannot recognize a token spent between other clients, even if it has already owned the token [23]. It has however been proven that an adversary can always recognize his previously-owned tokens if they are paid back to him [23].

Furthermore, it is assumed that authorities know the identities of their clients. It is expected that fraudsters cannot always be penalized within the confines of the transaction system. For example, dealing a corrective fine would require a convict to own enough tokens to pay. If a fine cannot be paid, corrective actions need to be taken in another way that does not involve tokens. Finding a fair way to correct fraud and penalize fraudsters was intentionally left out of scope.

V. IMPLEMENTATION

We prototyped the design described in Section IV. The prototype deliberately includes only the basic facets required to transact currency per Section IV. It was implemented on the Kotlin-IPv8 protocol stack. Kotlin-IPv8 provides utilities for peer-to-peer communication, such as UDP hole punching and public key cryptography. The specified sizes for public keys and signatures mentioned in Section IV-A were adapted from Kotlin-IPv8.

The *de facto* way of transferring large binary files via Kotlin-IPv8 is by using the EVA protocol [24]. EVA is an acknowledgement protocol for UDP that uses acknowledgement windows to guarantee packet delivery. We were confronted with EVA's limitations with regards to stability and throughput and thus opted to use our own (slightly) modified version. Our modified version LINK NAAR PR:

- Fixes a race condition that caused EVA to fail subsequent data transfers arbitrarily.
- Uses a faster and more compact encoding of lost packet numbers. HOEVEEL SNELLER?
- Allows encryption to be disabled.

Nevertheless, in Section VI-C we will demonstrate that we were unable to solve all problems with EVA that we encountered.

Furthermore, the implementation provides several scripts that were used for the performance analysis that will follow in Section VI. For details, please refer to Table I.

VI. PERFORMANCE ANALYSIS

We present an analysis of our prototype's performance. For a proper frame of reference, we place additional emphasis on the low-level functionality of the underlying Kotlin-IPv8 framework.

Experiments were performed on standard consumer electronics; a Lenovo Thinkpad L13 with an Intel i5 CPU operating at 2.11 GHz and 8 GB of DDR4 RAM. All experiments were performed 10 times; the variance of the results is depicted in most figures.

A. Cryptographic Verification

We measured the throughput of cryptographic signing and verification operations to ascertain performance bounds of Kotlin-IPv8 and thus the transaction protocol. The core idea is that by stripping the implementation of all other factors, the influence of cryptographic operations on an authority's throughput can be determined. All operations were performed with Ed25519 [25] using a Kotlin port of Libsodium⁴ that is also used by Kotlin-IPv8. The chosen parameters were identical to those used in Kotlin-IPv8.

Figure 3 shows throughput of the cryptographic operations required to sign and verify tokens in an online setting. In this setting, three signatures are to be verified: the authority's initial signature (147); the signature of the first recipient to the second; and the signature of the second back to the authority. The figure shows that throughput increases with the number of CPUs until a certain threshold. This limits scalability significantly. We suspect the diminishing increase to be due to resource sharing within Libsodium, although the exact reasons are unknown. Interestingly, the highest verification measurement of 22489 tokens per second, at $147+138+138 = 423$ bytes to verify per token, corresponds to a verification throughput of only 9.51 megabytes per second. To explain this, we measured signature verification for different data sizes.

Figure 4 shows the throughput of cryptographic signing and verification for varying data sizes on a single thread. Again we used ed25519 and Libsodium to mimic Kotlin-IPv8. It is apparent that larger file sizes are tremendously faster to verify than smaller. It is worth noting again that each token requires three separate signature operations on files that are all < 147 bytes.

⁴For Lazysodium, see <https://github.com/terl/lazysodium-java>.

TABLE I
FACETS OF PROTOTYPE AND ANALYSIS

	Prototype			Analysis						
Feature	Client	Verifier	Minting	Cryptography	Pipe	UDP (native)	UDP (Kotlin)	Kotlin-IPv8	EVA	Plots
LoC	10 (10%)	10 (10%)	10 (10%)	10 (10%)	10 (10%)	10 (10%)	10 (10%)	10 (10%)	10 (10%)	10 (10%)
Language	Kotlin	Kotlin	Kotlin	Kotlin	C	C	Kotlin	Kotlin	Kotlin	Python




Fig. 3. The benefits of parallelism diminish quickly. Throughput of *only* cryptographic verification and signing of tokens.

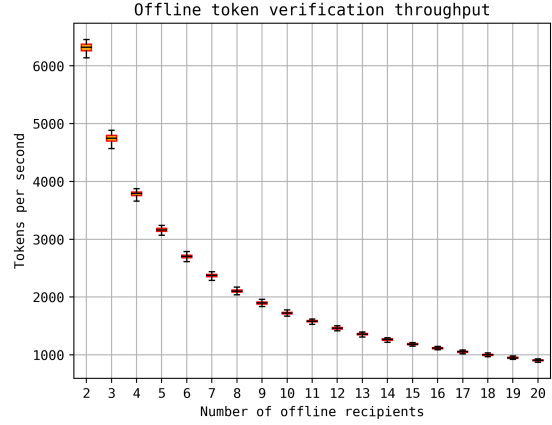


Fig. 5. Throughput of token verification declines for additional recipients.

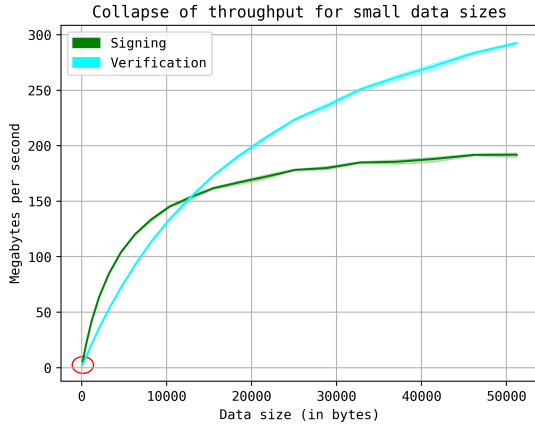


Fig. 4. Throughput of cryptographic signing and verification collapses for small data sizes. Operations on tokens are at most 147 bytes, marked by the red circle.

Figure 5 show that the impact of verification on throughput is also noticeable when tokens are transacted in an offline setting and require additional verification after the fact.

B. Data Transfer

We performed a range of data transfers with increasingly more high-level software. Measurements were performed on one machine, to eliminate the fluctuations of network latency. Figure 6 shows there are large drops in throughput when switching from a native program to the JVM; again when data packets are newly allocated instead of reusing the same

memory space; and again when switching to Kotlin-IPv8. However, the former two are unpreventable given the scope of this thesis.

Within the scope of this thesis are the measurements regarding Kotlin-IPv8 and EVA. Based on the results of Figure 4, we expect per-packet encryption to also be a bottleneck for Kotlin-IPv8’s throughput. The results for EVA are detailed in Section VI-C.




Fig. 6. Throughput of various data transfer methods.

C. EVA

Figure 7 shows EVA’s throughput for various configurations. We performed our measurements on a single machine to eliminate external noise. EVA’s main configurable parameters are the number of payload bytes per UDP packet (*block size*) and the number of blocks per acknowledgement window (*blocks per window*). We observed a noticeably low throughput for some configurations of EVA. From the figure it can be observed that the entire column of 256 blocks per window shows subpar throughput when encryption is *disabled*. Notice how there is a positive correlation between throughput and

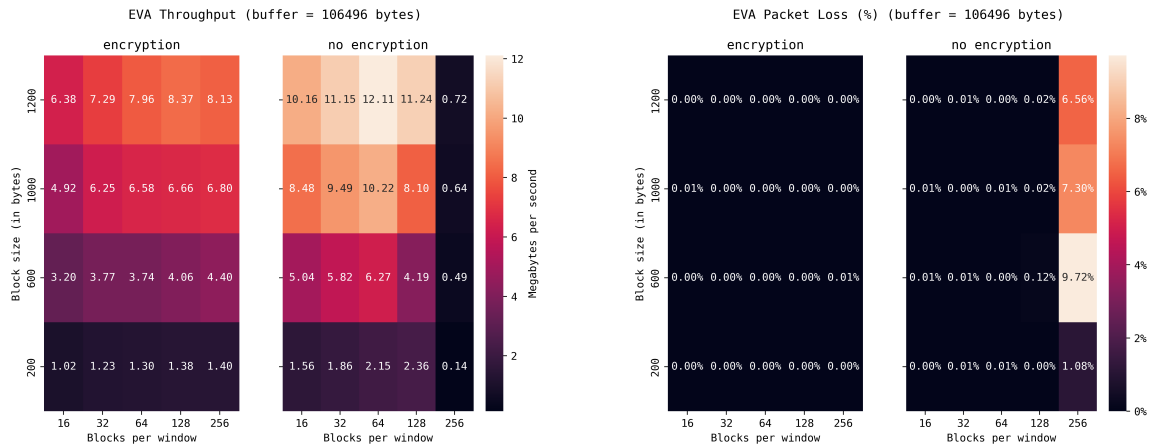


Fig. 7. Throughput (left) and packet loss (right) for various configurations of EVA. The UDP send and receive buffers were fixed at 106496 bytes.

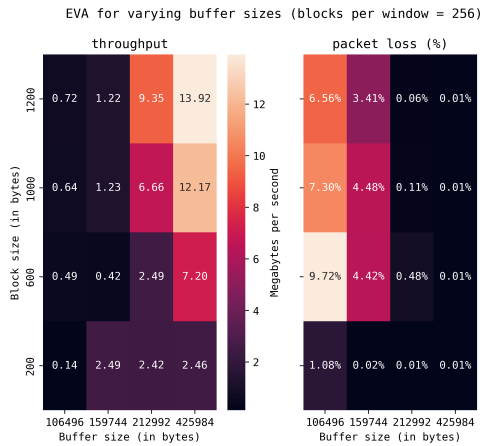


Fig. 8. Throughput of EVA for various UDP buffer sizes. The number of blocks per window was fixed at 256.

both block size and blocks per window, save for the last two columns of *no encryption*.

To explain the sudden drop in throughput, Figure 7 shows the measurements’ corresponding packet loss as well. From the figure it can be observed that packet loss is non-negligible only in the last two columns of *no encryption*. Our measurements were performed locally, so external routing devices cannot account for the anomaly. The fact that none of the *encryption* columns suffer from packet loss indicates that processing speed is not the bottleneck either; enabling encryption is strictly more intensive than not. We therefore hypothesized the drop to be caused by UDP buffers overflowing.

Figure 8 shows that this is indeed the case. Packet loss can be mitigated by increasing the UDP send and receive buffers. From inspection of the source code, we found that Kotlin-IPv8’s EVA implementation attempts to transmit an entire window’s worth of data at once. For example, at 256 blocks per window and blocks of 1200 bytes, this results in a series of consecutive transmissions amounting to $256 \cdot 1200 = 307.2\text{KB}$

excluding overhead. This transmission causes buffer overflows and thereby packet loss. We found that enabling encryption acts as an unintended form of *congestion control*. Due to reduced throughput (see Figure 4), the buffer does not fill up. We recommend improving EVA with a proper form of congestion control but have left this out of scope.

D. Tokens

MEET MET 40K tokens, dat is ongeveer 10MB aan payload.o CLIENT RECEIVE SPEED AUTHORITY VERIFY SPEED AUTHORITY SIGN SPEED

VII. CONCLUSION & FUTURE WORK

This thesis describes a token-based transaction protocol and its implementation on the Kotlin-IPv8 protocol stack. The protocol allows funds to be spent in an offline setting and guarantees retroactive fraud detection. During the performance analysis of our prototype we found that Kotlin-IPv8 in its current state is unable to properly fulfill the requirements for a European CBDC.

Further improvements to Kotlin-IPv8 are necessary; especially in the fields of data transfer and cryptography. The implemented prototype also requires improvements, not least with regards to anonymity, throughput, and decentralization.

REFERENCES

- [1] ECB, “Report on a digital euro,” ECB, Tech. Rep., 2020. [Online]. Available: https://www.ecb.europa.eu/pub/pdf/other/Report_on_a_digital_euro~4d7268b458.en.pdf
- [2] F. Panetta, “Public money for the digital era: towards a digital euro,” 2022, keynote speech by Fabio Panetta, Member of the Executive Board of the ECB. [Online]. Available: <https://www.ecb.europa.eu/press/key/date/2022/html/ecb.sp220516~454821f0e3.en.html>
- [3] ECB, “Call for expression of interest for digital euro front-end prototyping,” 2022, a call for expression of interest by ECB regarding a front-end prototype for a digital Euro. [Online]. Available: https://www.ecb.europa.eu/paym/digital_euro/investigation/profuse/shared/files/dedocs/ecb.dedocs220428.en.pdf
- [4] —, “Faq’s on the digital euro,” 2022, an FAQ by ECB for the digital Euro project. [Online]. Available: https://www.ecb.europa.eu/paym/digital_euro/faqs/html/ecb.faq_digital_euro.en.html

- [5] N. Jonker, L. Hernandez, R. de Vree, and P. Zwaan, "From cash to cards: how debit card payments overtook cash in the netherlands," DNB, Tech. Rep., 2018. [Online]. Available: https://www.dnb.nl/media/kx1akmnb/201802_nr_1_-2018-_from_cash_to_cards_how_debit_card_payments_overtook_cash_in_the_netherlands.pdf
- [6] N. Jonker and P. Zwaan, "Betalen aan de kassa 2020," DNB, Tech. Rep., 2020. [Online]. Available: https://www.dnb.nl/media/e34bo5zu/betalen_kassa_2020.pdf
- [7] S. Riksbank, "Pandemic hastening development towards digital payments," 2021, an article published by Sveriges Riksbank about the declining usage of cash in Sweden and other countries. [Online]. Available: <https://www.riksbank.se/en-gb/payments--cash/payments-in-sweden/payments-report-2021/1.-trends-on-the-payment-market/pandemic-hastening-development-towards-digital-payments/cash-used-less-and-less-frequently-in-sweden-and-abroad/>
- [8] ECB, "The paradox of banknotes: understanding the demand for cash beyond transactional use," *ECB Economic Bulletin, Issue 2/2021*, 2021. [Online]. Available: https://www.ecb.europa.eu/pub/economic-bulletin/articles/2021/html/ecb.ebart202102_03~58cc4e1b97.en.html
- [9] BIS, "G7 working group on stablecoins," BIS, Tech. Rep., 2019. [Online]. Available: <https://www.bis.org/cpmi/publ/d187.pdf>
- [10] Office of Science and Technology Policy, "Technical evaluation for a u.s. central bank digital currency system," Office of Science and Technology Policy, Tech. Rep., 2022. [Online]. Available: <https://www.whitehouse.gov/wp-content/uploads/2022/09/09-2022-Technical-Evaluation-US-CBDC-System.pdf>
- [11] Working Group on E-CNY Research and Development of the People's Bank of China, "Progress of research & development of e-cny in china," Working Group on E-CNY Research and Development of the People's Bank of China, Tech. Rep., 2021. [Online]. Available: <http://www.pbc.gov.cn/en/3688110/3688172/4157443/4293696/2021071614584691871.pdf>
- [12] House of Lords, "Central bank digital currencies: a solution in search of a problem?" 2022. [Online]. Available: <https://publications.parliament.uk/pa/ld5802/ldselect/ldconaf/131/131.pdf>
- [13] National Commission on Informatics and Liberty, "Digital euro: what is at stake for privacy and personal data protection?" 2022. [Online]. Available: <https://www.cnil.fr/en/digital-euro-what-stake-privacy-and-personal-data-protection>
- [14] S. Nakamoto, "Bitcoin: A peer-to-peer electronic cash system," *Decentralized Business Review*, p. 21260, 2008.
- [15] D. Chaum, "Blind signatures for untraceable payments," in *Advances in cryptography*. Springer, 1983, pp. 199–203.
- [16] S. Canard, D. Pointcheval, O. Sanders, and J. Traoré, "Divisible e-cash made practical," in *IACR International Workshop on Public Key Cryptography*. Springer, 2015, pp. 77–100.
- [17] T. Okamoto and K. Ohta, "Disposable zero-knowledge authentications and their applications to untraceable electronic cash," in *Conference on the Theory and Application of Cryptology*. Springer, 1989, pp. 481–496.
- [18] M. Stadler, J. Piveteau, and J. Camenisch, "Fair blind signatures," in *International Conference on the Theory and Applications of Cryptographic Techniques*. Springer, 1995, pp. 209–219.
- [19] R. Blokzijl, "Eurotoken," Master's thesis, Delft University of Technology, 2021. [Online]. Available: <http://resolver.tudelft.nl/uuid:132faae8-6883-454f-a8ce-94735340dce9>
- [20] V. Hajric and K. Greifeld, "Bitcoin went mainstream in 2021. it's just as volatile as ever," Bloomberg, 2021. [Online]. Available: <https://www.bloomberg.com/graphics/2021-bitcoin-volatility/>
- [21] D. Bullmann, J. Klemm, and A. Pinna, "In search for stability in crypto-assets: are stablecoins the solution?" ECB, Tech. Rep., 2019. [Online]. Available: <https://www.ecb.europa.eu/pub/pdf/scpops/ecb.op230%7Ed57946be3b.en.pdf>
- [22] S. Canard, A. Gouget, and J. Traoré, "Improvement of efficiency in (unconditional) anonymous transferable e-cash," in *International Conference on Financial Cryptography and Data Security*. Springer, 2008, pp. 202–214.
- [23] S. Canard and A. Gouget, "Anonymity in transferable e-cash," in *International Conference on Applied Cryptography and Network Security*. Springer, 2008, pp. 207–223.
- [24] J. Bambacht and J. Pouwelse, "Web3: A decentralized societal infrastructure for identity, trust, money, and data," Master's thesis, Delft University of Technology, 2022.
- [25] D. Bernstein, N. Duif, T. Lange, P. Schwabe, and B. Yang, "High-speed high-security signatures," *Journal of cryptographic engineering*, vol. 2, no. 2, pp. 77–89, 2012.

ACKNOWLEDGMENT

The preferred spelling of the word "acknowledgment" in America is without an "e" after the "g". Avoid the stilted expression "one of us (R. B. G.) thanks ...". Instead, try "R. B. G. thanks...". Put sponsor acknowledgments in the unnumbered footnote on the first page.

IEEE conference templates contain guidance text for composing and formatting conference papers. Please ensure that all template text is removed from your conference paper prior to submission to the conference. Failure to remove the template text from your paper may result in your paper not being published.